

Approximating a Multi-Grid Solver

Valentin Le Fèvre Leonardo Bautista-Gomez Marc Casas
`valentin.le-fevre@ens-lyon.fr`

ENS de Lyon, Barcelona Supercomputing Center

April 17, 2018

Introduction

Approximate computing: trade-off between **accuracy of result** and **execution time**.

Approximate computing: trade-off between **accuracy of result** and **execution time**.

- Precision of a floating-point value
- No exact result exists (search query...)
- ...

Approximate computing: trade-off between **accuracy of result** and **execution time**.

- Precision of a floating-point value
- No exact result exists (search query...)
- ...
- Skip steps in loops
- Branching to avoid useless computations
- Faulty hardware (fast adders...)
- ...

Introduction

- Multi-Grid (MG) solvers [3]: iterative solvers with different level of coarseness: number of evaluation points.
Faster than classical method and scales well.

4096 points

1024 points

256 points

64 points




Figure: Example of cycle: each blue point represents one iteration of an iterative method, while the height corresponds to the coarseness of the grid. Red is “exact” solve.

Introduction

- Multi-Grid (MG) solvers [3]: iterative solvers with different level of coarseness: number of evaluation points.
Faster than classical method and scales well.

4096 points

1024 points

256 points

64 points




Figure: Example of cycle: each blue point represents one iteration of an iterative method, while the height corresponds to the coarseness of the grid. Red is “exact” solve.

- Accuracy of result (*relative residual norm*) is limited by the hardware.
- We do not aim the same accuracy when using it as a conditioner or a solver.

- 1 The UP-cycle
- 2 Bitwidth, performance and accuracy
- 3 Conclusion

- 1 The UP-cycle
- 2 Bitwidth, performance and accuracy
- 3 Conclusion

First idea: add more iterations at each level or more complex cycles.




Figure: Relative residual norm and execution time as function of number of cycles.

First idea: add more iterations at each level or more complex cycles.




Figure: Relative residual norm as function of execution time.

Level	Matrix size	Non-zero elements	Relax (down)	Relax (up)	Restriction	Interpolation
1	512,000	4,042,520	20 ms	20 ms	15 ms	-
2	256,000	6,475,239	20 ms	25 ms	12 ms	4 ms
3	58,893	2,000,513	8 ms	8 ms	3 ms	2 ms
4	14,285	788,509	2 ms	2 ms	1 ms	0.7 ms
5	4,238	386,333	1 ms	1 ms	0.5 ms	0.2 ms
6	609	53,493	< 0.1 ms	< 0.1 ms	< 0.1 ms	< 0.1 ms
7	69	2,873	< 0.1 ms	< 0.1 ms	< 0.1 ms	< 0.1 ms
8	2	4	< 0.1 ms	-	-	< 0.1 ms

Table: Time breakdown of a V-cycle with $\alpha = 1$.

⇒ Relaxations represent $\approx 66\%$ of the total cost of a V-cycle.

The U_P -cycle

After several tries: the U_P -cycle.

We do relaxations only when going up in the V-cycle.









Figure: Blue: relaxation. Red: exact solve. Black: nothing.


Results




Results



(a) $3 \times 3 \times 3$



(b) $6 \times 6 \times 1$



(c) $4 \times 4 \times 4$

Overall, between **7%** and **28%** of improvement for reaching max accuracy on our tests.

- 1 The UP-cycle
- 2 Bitwidth, performance and accuracy
- 3 Conclusion

Impact of bitwidth

- The bitwidth is a hardware limitation: we can't have results accurate to 2^{-1000} using double floating-point representation.
- However, using a small bitwidth makes computations faster and more energy-efficient [2].

Impact of bitwidth

- The bitwidth is a hardware limitation: we can't have results accurate to 2^{-1000} using double floating-point representation.
- However, using a small bitwidth makes computations faster and more energy-efficient [2].
- We rewrite the MG algorithm: one version using only single-precision floating-points and one version with the relaxation algorithm using MPFR variables (arbitrary precision) [1].

Impact of bitwidth

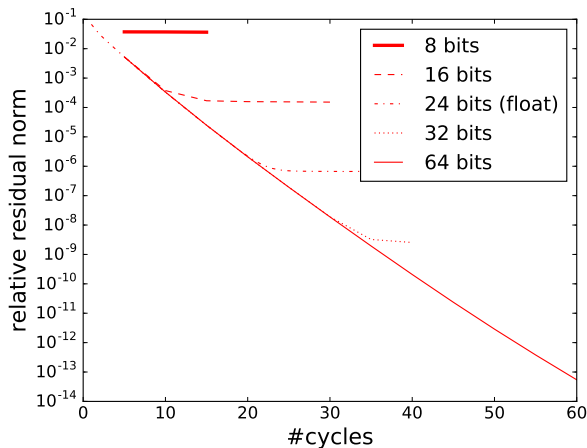


Figure: Accuracy for different number of **mantissa** bits.

Impact of bitwidth

- The bitwidth is a hardware limitation: we can't have results accurate to 2^{-1000} using double floating-point representation.
- However, using a small bitwidth makes computations faster and more energy-efficient [2].
- We rewrite the MG algorithm: one version using only single-precision floating-points and one version with the relaxation algorithm using MPFR variables (arbitrary precision) [1].
- Conclusion: using small bitwidths does not change the convergence rate (until late).

t a threshold, $\text{UPDATE}(b)$ a function which returns an integer greater than b .

- ① $b \leftarrow 64$.
- ② **While** $\text{nb_iters} < \text{max_iter}$ **and** $\text{rel_res_norm} > \text{tolerance}$
 - ① Do a cycle at precision b .
 - ② Compute new_rel_res_norm .
 - ③ $\text{rel_res_norm} \leftarrow \text{new_rel_res_norm}$.
 - ④ $\text{nb_iters} \leftarrow \text{nb_iters} + 1$.

t a threshold, $\text{UPDATE}(b)$ a function which returns an integer greater than b .

- ① $b \leftarrow 16$.
- ② **While** $\text{nb_iters} < \text{max_iter}$ **and** $\text{rel_res_norm} > \text{tolerance}$
 - ① Do a cycle at precision b .
 - ② Compute new_rel_res_norm .
 - ③ **If** $\text{new_rel_res_norm} > t \times \text{rel_res_norm}$ **Then** $b \leftarrow \text{UPDATE}(b)$.
 - ④ $\text{rel_res_norm} \leftarrow \text{new_rel_res_norm}$.
 - ⑤ $\text{nb_iters} \leftarrow \text{nb_iters} + 1$.

Algorithm

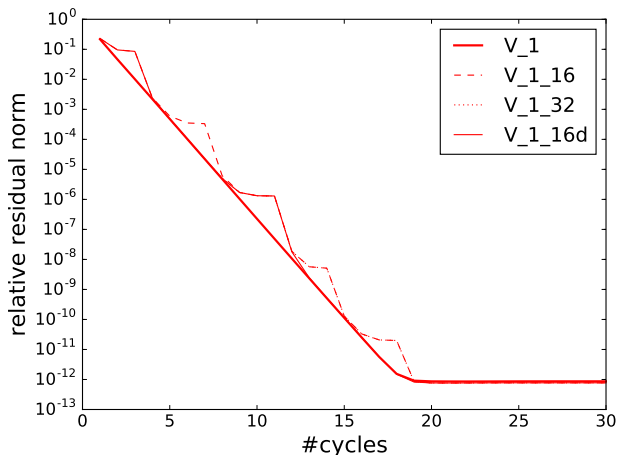


Figure: Accuracy of adaptive algorithms compared to the original double-precision with a precision threshold of 0.8.

How to estimate the benefits in term of execution time ?

How to estimate the benefits in term of execution time ?

$$Time(n, b) = a \cdot n^3 \cdot b^\alpha + c$$

- n : size of the problem (we worked on 3D grids so n^3 for the size of the matrix).
- b : number of mantissa bits.
- a, α, c : constants to determine.

We found $\alpha \approx 0.3$ (sublinear...) using single-precision and double-precision codes.

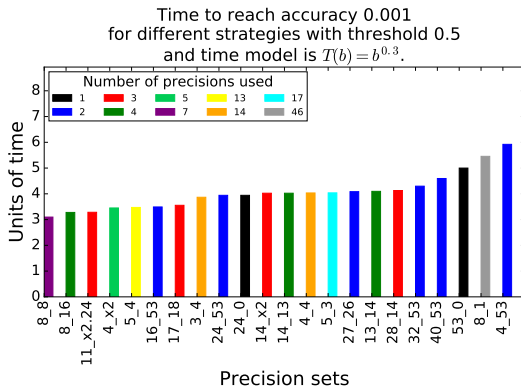


Figure: Cost of the MG solver considering several different dynamic precision scenarios to reach accuracy 10^{-3} .

GPU compared to double-precision: **34%** improvement.




Figure: Cost of the MG solver considering several different dynamic precision scenarios to reach accuracy 10^{-7} .

GPU compared to double-precision: **23%** improvement.




Figure: Cost of the MG solver considering several different dynamic precision scenarios to reach accuracy 10^{-15} .

GPU compared to double-precision: **9%** improvement.

- 1 The UP-cycle
- 2 Bitwidth, performance and accuracy
- 3 Conclusion**

Conclusion (1)

Final results:

- A new cycle shape that tends to converge faster: the UP-cycle.
- A new algorithm for *any* MG cycle shape that reduces execution time and energy consumption.
- Up to 30% expected improvement on a GPU with half-precision/single-precision/double-precision available by mixing UP-cycle and changing bitwidths.
- At least 15% expected improvement for reaching maximum accuracy compared to previously.

Conclusion (2)

Future ideas:

- Model (or measure) the gains in energy consumption instead of execution time.
- Change precision **inside** a cycle?
- Link to silent data corruption: what if the environment forces us to work at 10^{-x} as max accuracy because of bitflips?

Conclusion (2)

Future ideas:

- Model (or measure) the gains in energy consumption instead of execution time.
- Change precision **inside** a cycle?
- Link to silent data corruption: what if the environment forces us to work at 10^{-x} as max accuracy because of bitflips?

Thank you for your attention. Any question?

References



Laurent Fousse, Guillaume Hanrot, Vincent Lefèvre, Patrick Pélissier, and Paul Zimmermann.

Mpfr: A multiple-precision binary floating-point library with correct rounding.

ACM Trans. Math. Softw., 33(2), June 2007.



Gokul Govindu, Ling Zhuo, Seonil Choi, Padma Gundala, and Viktor K Prasanna.

Area and power performance analysis of a floating-point based application on FPGAs.

In Seventh Annual Workshop on High Performance Embedded Computing.




W. Hackbusch.


Multi-Grid Algorithms, pages 133–160.

Springer Berlin Heidelberg, Berlin, Heidelberg, 1991.

Grids



Grids



Grids

